

Database Management Systems

Reagan Shirk

August 31, 2020

View of Data

- You have a logical level and a physical level of data
- The physical level is telling the management system how you want your data to physically be stored

Relational Model

- A collection of tables...?
 - I didn't totally catch what she was saying here but it's in Chapter 2 of the book and I'll go look and try to elaborate
 - Oh I think it's a collection of similar data, for example she has a table of instructors and a table of departments
- When talking about the logical schema, you give the datatype for each column in the data
 - i.e. if you have an ID column and a Name column, the ID would be an int and the Name would be a string

Data Definition Language (DDL)

- Language used for defining the database schema
- Relational database uses SQL

```
create table instructor(  
    ID          char(5),  
    name        varchar(20),  
    dept_name   varchar(20),  
    salary      numeric(8, 2))  
)
```

- If you're curious, the difference between `char` and `varchar` is that `char` is of fixed length and `varchar` is of variable length
- The DDL compiler creates a set of table templates stored in a data dictionary (this data dictionary is metadata - data about data)
- There are integrity constraints; primary key and referential integrity
 - For example, the `dept_name` value in any `instructor` tuple must appear in the `department` relation

SQL (Chapters 3, 4, 5)

- SQL is a widely used non-procedural language
- Applications generally access databases through either:
 - language extensions that allow embedded SQL
 - interface which allows for SQL queries to be sent to a database

Database Design

- A longwinded way of saying that you don't want redundancy in your data. Why would the department budget need to be in a datatable of professor information? It wouldn't.

Design Approaches

Entity Relationship Model

- A collection of *entities* and *relationships* (I'm shocked)
 - An entity is a thing or object in the enterprise that is distinguishable from other objects, it is described by a set of attributes
 - A relationship is an association among several entities
 - * An instructor is a member of a department, that is their entity-relationship model
- This is **not an implementation model**, it is just for design
 - An implementation model is something that has been/can be implemented

Storage Management

- A storage manager is a program module that links the low-level data and the application programs/queries submitted to the system
- It is responsible for:
 - interaction with the file manager
 - Efficient storing, retrieving, and updating
- Issues:
 - Storage access
 - File organization
 - Indexing and hashing

Query Processing

- You have your query and your DBMS
 - You submit your query and it goes into the query processor to be parsed and translated
 - I got distracted. Looks like it goes into an optimizer after doing something with algebra? I'll have to come rewatch this part of lecture
 - After it gets optimized, the execution is planned and evaluation is performed, then you get your output

Transaction Management

- What if the system fails, what if more than one user is updating the same data at the same time?
- This is when transaction management comes into play
- Something about a transaction having a beginning and an end
- You can recover a database if the transaction fails in the middle
- A transaction is a collection of operations that performs a logical function
- Transaction management ensures that the database remains consistent and correct despite system and transaction failures
- Concurrency control controls the interaction among concurrent transactions to ensure consistency

Database Administrator (a person, not a code thing)

- a person who has central control over the system. They take care of:
 - schema definitions (logical design)
 - storage structure and access-method definitions (physical design)
 - schema and physical-organization modification
 - granting authorization for access
 - maintenance
 - backing up the database
 - ensuring enough free disk space
 - monitoring jobs running and performance

Database Architecture

- The architecture of a database system is influenced by the computer system that it's running on
- You got:
 - centralized
 - client server
 - parallel (multi-processor)
 - distributed

History Time

- 50's and 60's - AKA the stone age
 - Data processing using magnetic tapes for storage - these only provided sequential access
 - Punch cards used for input
- 60's and 70's - still the stone age
 - Hard disks (had to read that twice) allowed for direct access to data
 - Network and heirarchical data models were used frequently
 - Ted Codd (mathematician from England who worked for IBM) defines the relational data model (the not implementation model we talked about earlier)
 - * Won the ACM turing award in 1981 for being a badass
 - * He did something with set operations like union and intersection

When to not use a DBMS

- Woah our DBMS professor isn't telling us that we need to use one all the time? I'm shook
- Don't use one when theres:
 - too much cost associated due to
 - * overhead of concurrency control, recovery, integrity, etc
 - * High investment in hardware, software, and training
 - * generality
 - additional problems may pop up, like:
 - * database is not properly designed
 - * database application is not properly implemented
 - Could be more desireable for a regular file system when:
 - * db and application are simple and well defined, not expected to change
 - * time critical applications
 - * something I missed